EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Advanced Database Systems

Introduction

Strategic project of TBU in Zlín, reg. no. CZ.02.2.69/0.0/0.0/16_015/0002204

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Content

- Current trends in data management & computing
- Big Data
- Relational vs. NoSQL databases
  - the value of relational databases
  - new requirements and NoSQL features
  - flexible data models
- Types of NoSQL databases
  - key-value stores, document databases, column-family databases, graph databases
  - principles and examples

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Current Trends: Big Data

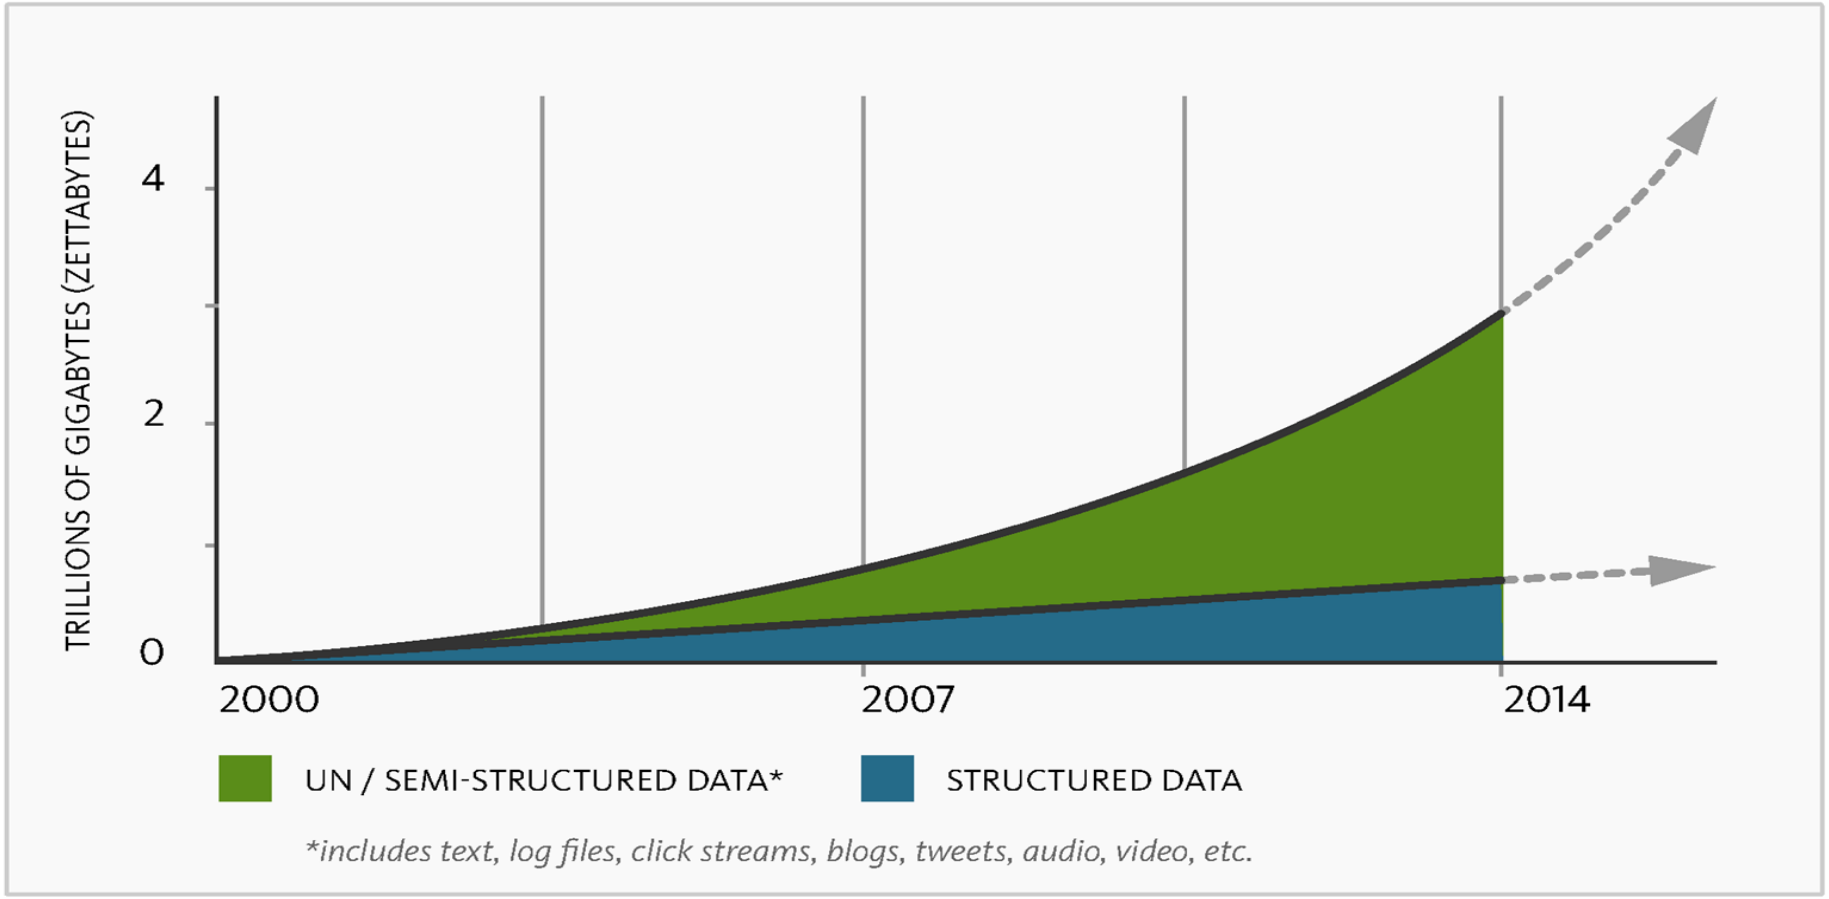

source: http://www.couchbase.com/sites/default/files/uploads/all/whitepapers/NoSQL-Whitepaper.pdf

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Current Trends: Big Users

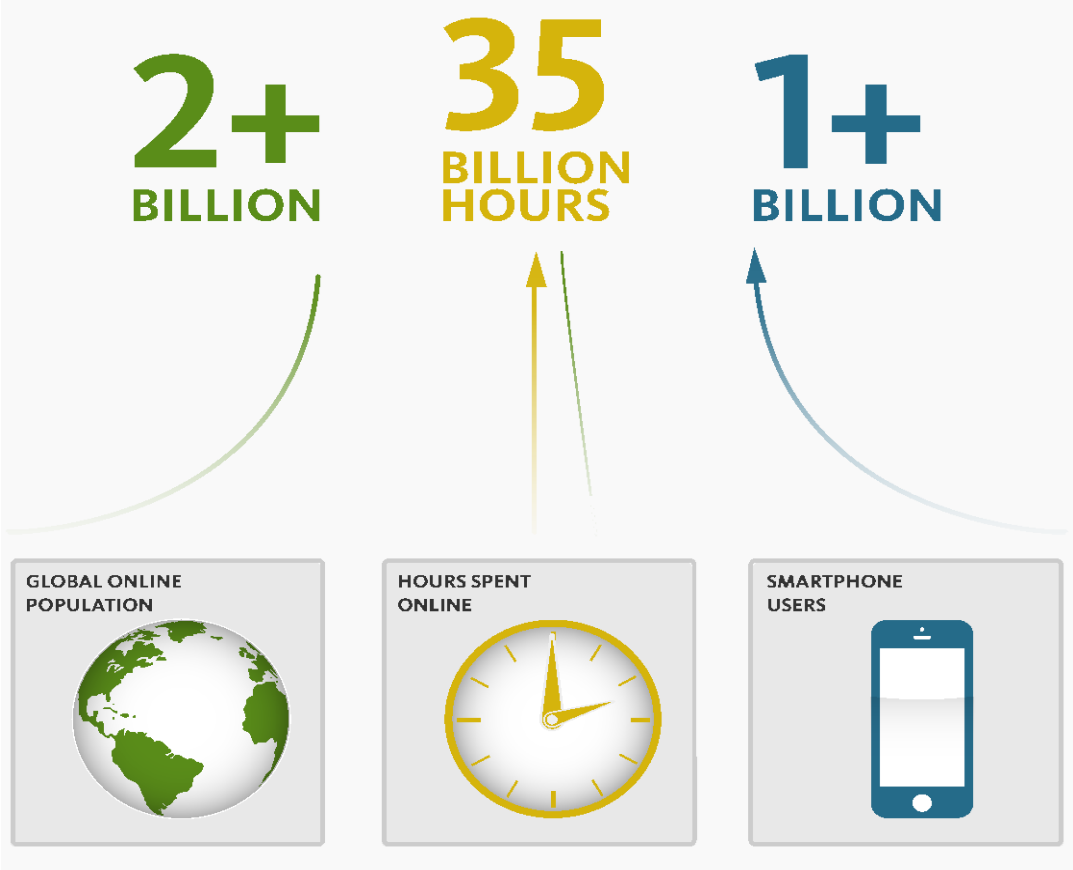

Zdenka Prokopova
TBU in Zlín

source: http://www.couchbase.com/sites/default/files/uploads/all/whitepapers/NoSQL-Whitepaper.pdf

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Current Trends: Cloud Computing

source: http://www.profitbricks.com/what-is-iaas

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Big Data

"Big data is high volume, high velocity, and/or high variety information assets that require new forms of processing to enable enhanced decision making, insight discovery and process optimization." (Gartner, 2012)

Zdenka Prokopova
TBU in Zlín

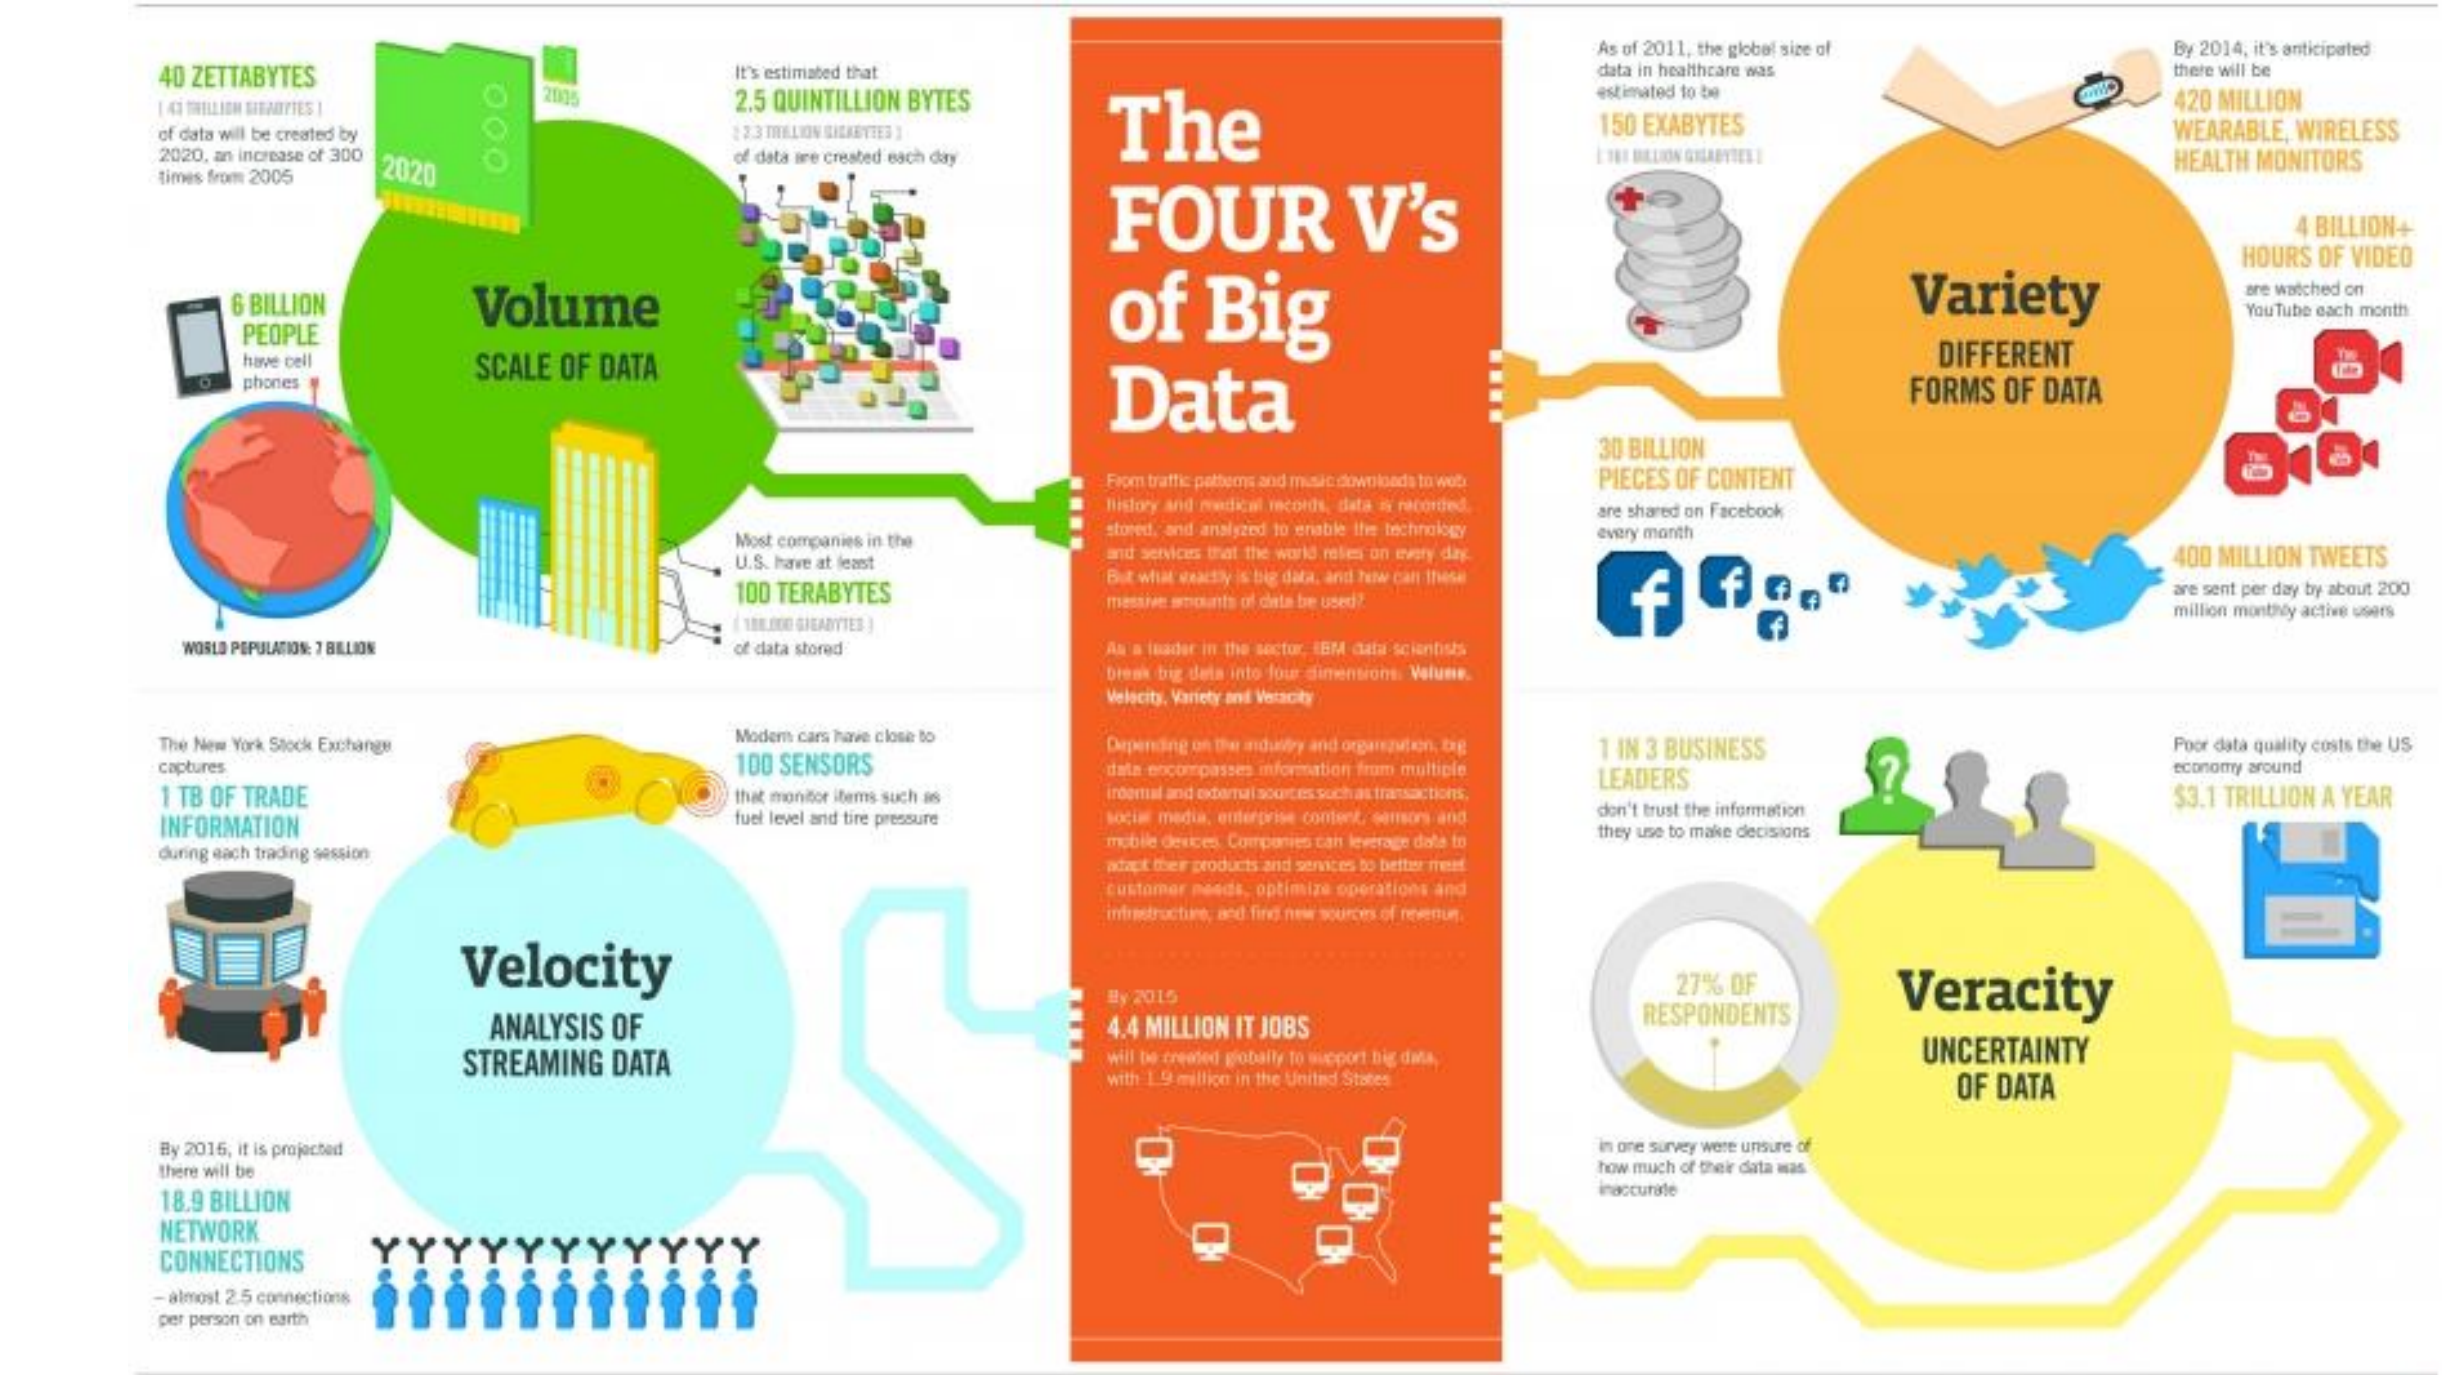

source: http://www.ibmbigdatahub.com/infographic/four-vs-big-data

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Data volume is increasing exponentially

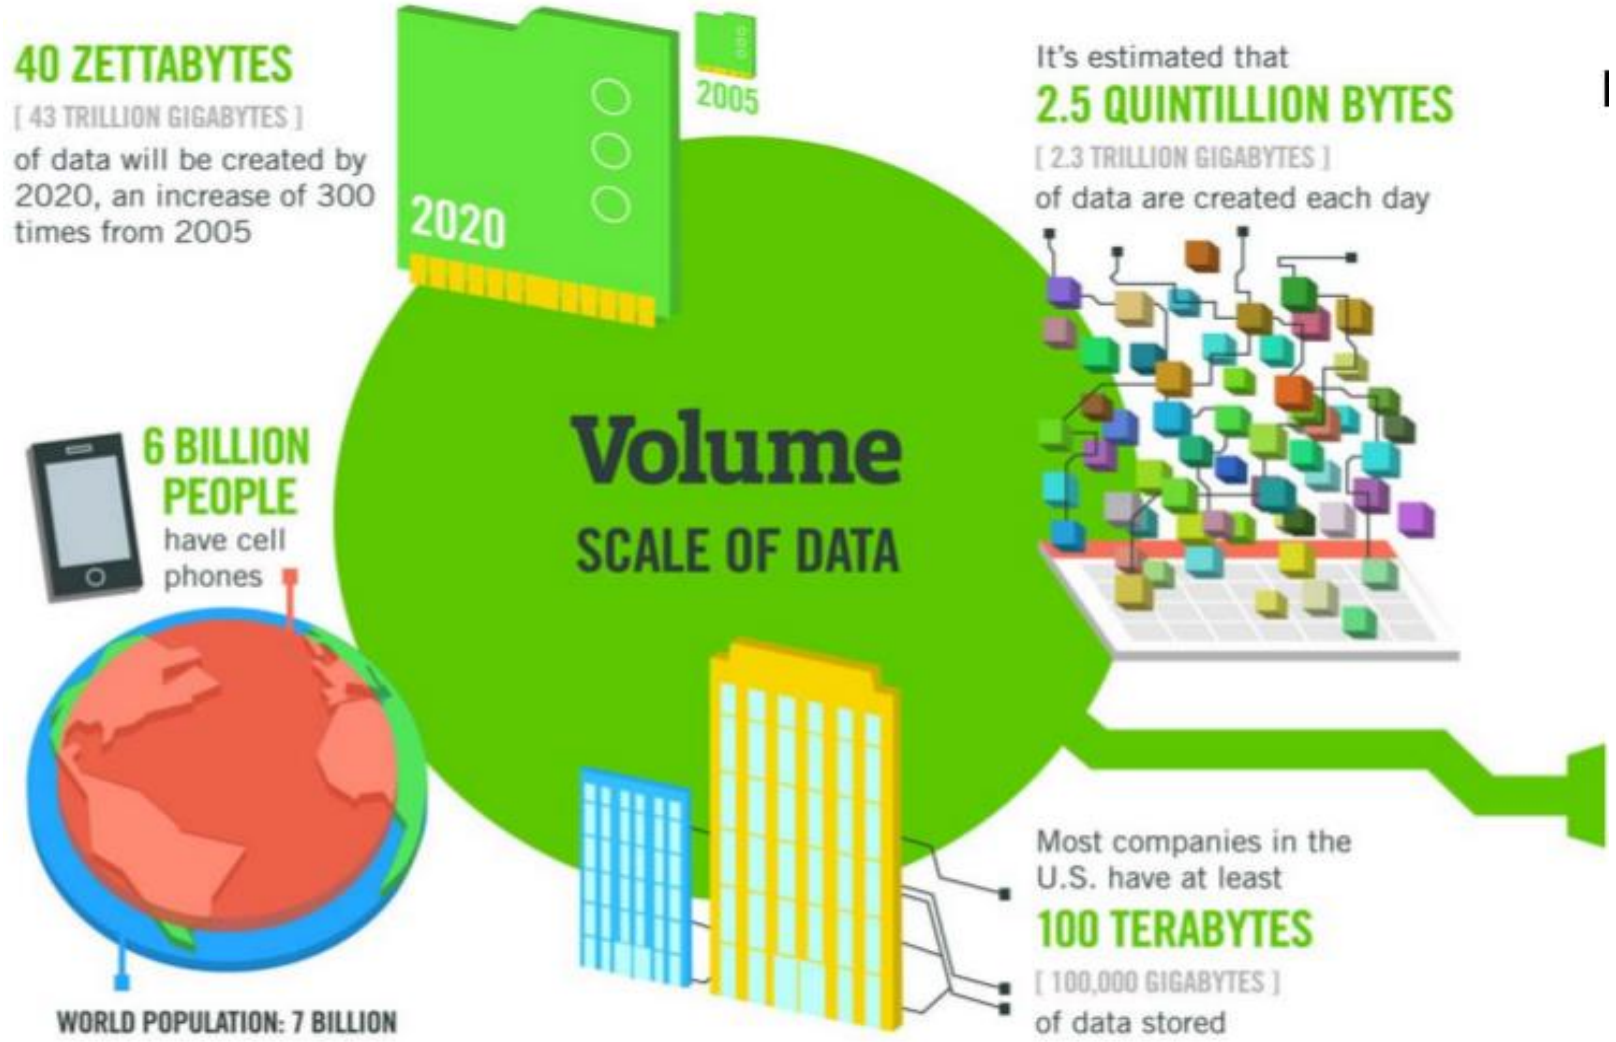

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Various data types, formats and structures

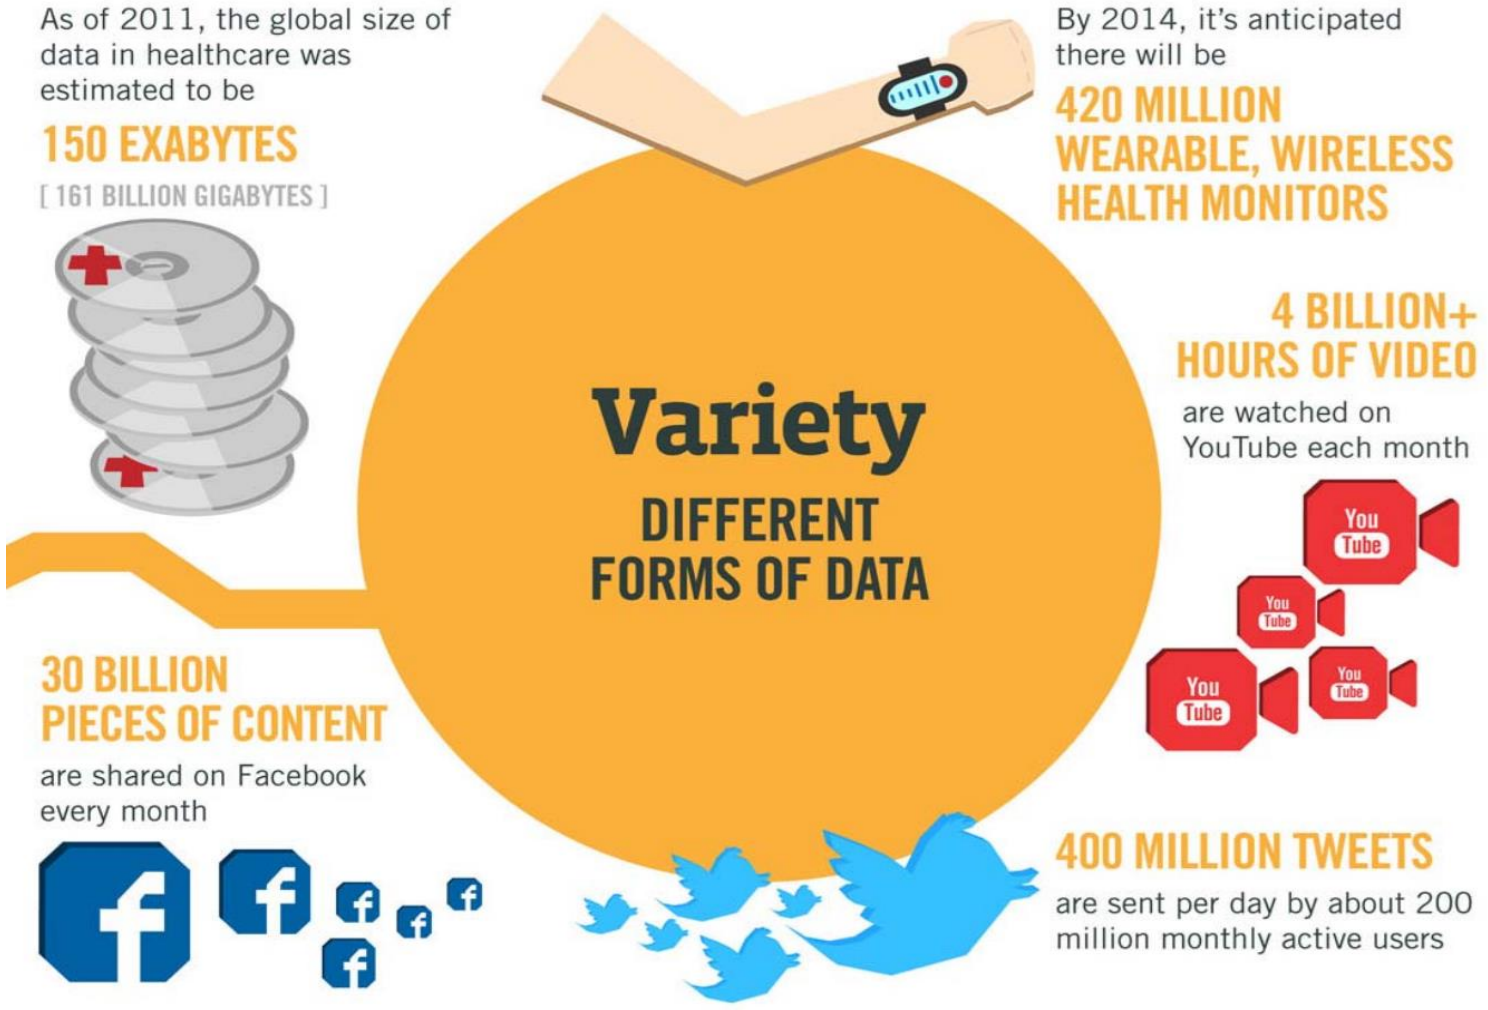

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Data is being generated fast and need to be processed fast

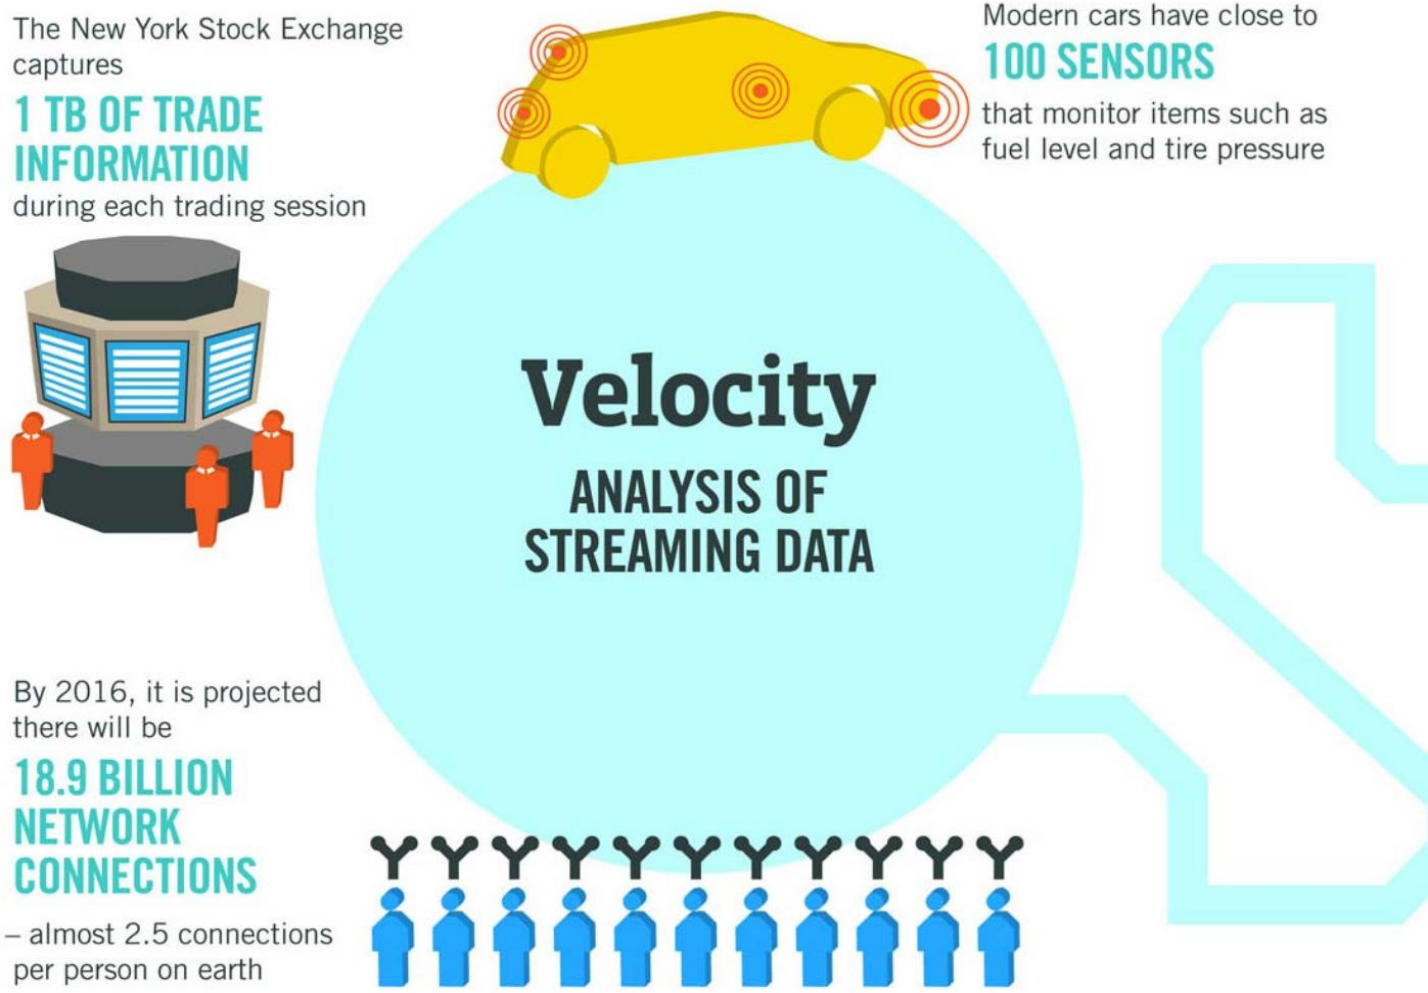

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Processing (Big) Data

- OLTP: Online Transaction Processing (DBMSs)
  - Database applications
  - Storing, querying, multi-user access
- OLAP: Online Analytical Processing (Warehousing)
  - Answer multi-dimensional analytical queries
  - Financial/marketing reporting, budgeting, forecasting, …
- RTAP: Real-Time Analytic Processing (Big Data Architecture & Technology)
  - Data gathered & processed in real-time (streaming)
  - Real-time and history data combined

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Technologies for Big Data

- Distributed file systems (GFS, HDFS, etc.)
- MapReduce
  - and other models for distributed programming
- NoSQL databases
- Grid computing, cloud computing
- Large-scale machine learning

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Relational Database Management Systems

- RDBMS are predominant database technologies
  - first defined in 1970 by Edgar Codd of IBM's Research Lab
- Data modeled as relations (tables)
  - object = tuple of attribute values
  - tables contain objects of the same type
  - tables interconnected via foreign keys
- Relational calculus, SQL query language

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

The Value of Relational Databases

- A (mostly) standard data model
- Many well developed technologies
  - physical organization of the data
  - search indexes: B^+-Trees, hash indexes
  - query optimization, search operator implementations
- Good concurrency control (ACID)
  - transactions: atomicity, consistency, isolation, durability
- Many reliable integration mechanisms
  - “shared database integration” of applications

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

New Requirements on Data Management

Trends

- **Volume** of data
- **Cloud** comp. (IaaS)
- **Velocity** of data
- **Big** users
- **Variety** of data

Requirements

- Real data **scalability**
  - massive database distribution
  - dynamic resource management
  - horizontally scaling systems
- Frequent **update** operations
- Massive **read** throughput
- **Flexible** database schema

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

NoSQL Databases

- What is "NoSQL"?
  - term used in late 90s for a different type of technology: Carlo Strozzi: http://www.strozzi.it/cgi-bin/CSA/tw7/I/en_US/NoSQL/
  - "Not Only SQL"?
    - but many RDBMS are also "not just SQL"

"NoSQL is an accidental term with no precise definition"

- first used at an informal meetup in 2009 in San Francisco (presentations from Voldemort, Cassandra, Dynomite, HBase, Hypertable, CouchDB, and MongoDB)

[Sadalage & Fowler: NoSQL Distilled, 2012]

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

NoSQL Databases

- NoSQL: Database technologies that are (mostly):
  - Not using the relational model (nor the SQL language)
  - Designed to run on large clusters (horizontally scalable)
  - No schema - fields can be freely added to any record
  - Open source
  - Based on the needs of 21st century web estates

Zdenka Prokopova
TBU in Zlín

[Sadalage & Fowler: NoSQL Distilled, 2012]

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

NoSQL Databases

- Other characteristics (often true):
  - easy replication support (fault-tolerance, query efficiency)
  - simple API
  - eventually consistent (not ACID)

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Assumptions about Data and Usage

	RDBMS	NoSQL
integrity	is mission-critical	OK as long as most data is correct
data format	consistent, well-defined	unknown or inconsistent
data	is of long-term value	is expected to be replaced
growth	predictable, linear growth	unpredictable growth (exponential?)
querying	non-programmers writing queries	only programmers writing queries
fault tolerance	regular backups	automatic data replication
distribution	access through master server	data sharding (partitioning)

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

The End of Relational Databases?

- Relational databases are not going away
- Many projects would use RDBMS also because of:
  - maturity/stability,
  - available support
  - familiarity

We should see RDBMS as one option for data storage

Polyglot persistence – using different data stores in different circumstances

Zdenka Prokopova
TBU in Zlín

[Sadalage & Fowler: NoSQL Distilled, 2012]

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Data Model: Aggregates

- The model by which the database organizes data
- Each NoSQL type has a different data model
  - Key-value, document, column-family, graph
  - First three are oriented on aggregates

Aggregate

- A data unit with a complex structure
  - Not simply a tuple like in RDBMS
- An aggregate is a collection of related objects that we wish to treat as a unit
  - unit for data manipulation and management of consistency

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Example: UML Model of an e-shop

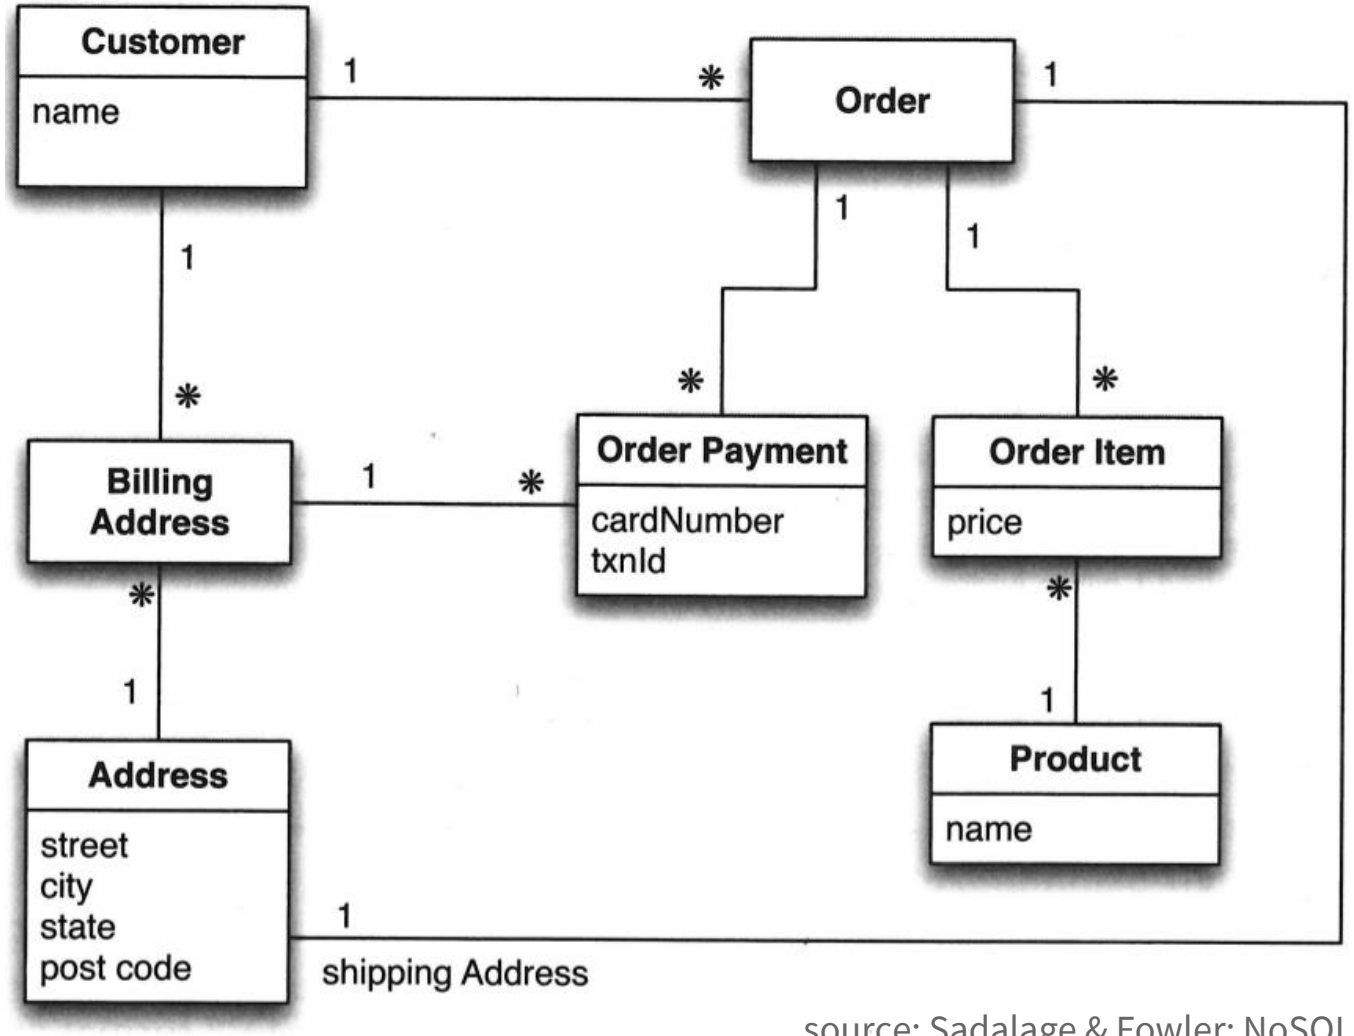

Zdenka Prokopova
TBU in Zlín

source: Sadalage & Fowler: NoSQL Distilled, 2012

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Example: Relational Model

Customer

Id	Name
1	Martin

Orders

Id	CustomerId	ShippingAddressId
99	1	77

Product

Id	Name
27	NoSQL Distilled

BillingAddress

Id	CustomerId	AddressId
55	1	77

OrderItem

Id	OrderId	ProductId	Price
100	99	27	32.45

Address

Id	City
77	Chicago

OrderPayment

Id	OrderId	CardNumber	BillingAddressId	txnId
33	99	1000-1000	55	abelif879rft

source: Sadalage & Fowler: NoSQL Distilled, 2012

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Relational Model: Aggregate Ignorant

- Relational databases are aggregate-ignorant
  - It is not a bad thing, it is a feature
  - Allows to easily look at the data in different ways
  - Best choice when there is no primary structure for data manipulation

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Example: NoSQL Solution

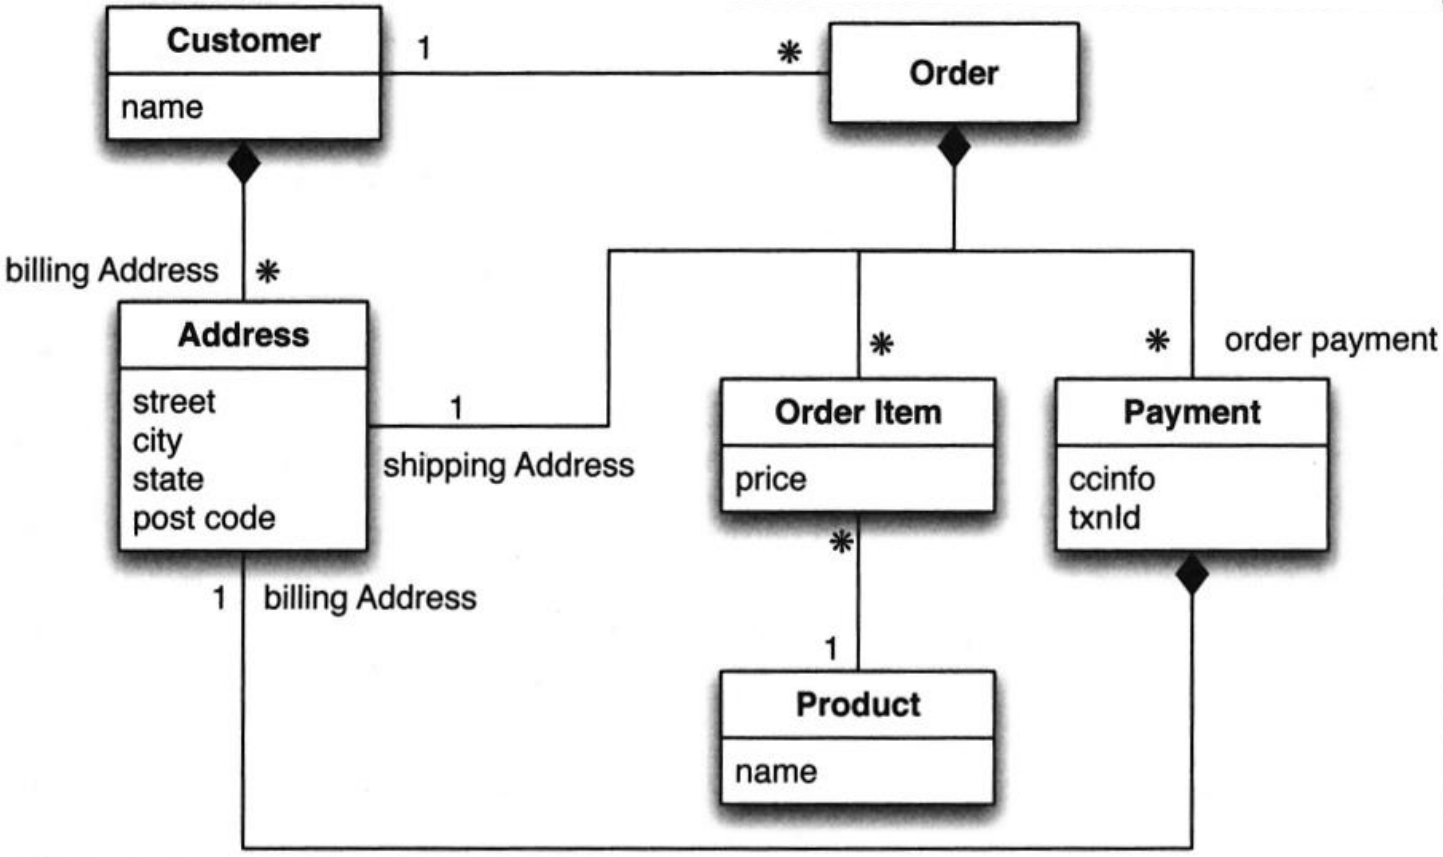


```
// in customers
{
"id":1,
"name":"Martin",
"billingAddress":[{"city":"Chicago"}]
}

// in orders
{
"id":99,
"customerId":1,
"orderItems":[
  {
  "productId":27,
  "price": 32.45,
  "productName": "NoSQL Distilled"
    }
  ],
"shippingAddress":[{"city":"Chicago"}]
"orderPayment":[
  {
    "ccinfo":"1000-1000-1000-1000",
    "txnId":"abelif879rft",
    "billingAddress": {"city": "Chicago"}
  }
  ],
}
```

source: Sadalage & Fowler: NoSQL Distilled, 2012

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

NoSQL Databases: Aggregate-oriented

- NoSQL databases are typically either:
  - schemaless (with implicit schema maintained by application)
  - or aggregate-oriented (more or less explicit schema)

Aggregate-oriented:

- There is no general strategy to set aggregate boundaries
- Aggregates give the database information about which bits of data will be manipulated together
  - Which should be stored on the same node

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Aggregate-oriented

Aggregates

- Helps greatly with running on a cluster of nodes
  - Minimize the number of nodes accessed during a search

- Impact on concurrency control:
  - NoSQL databases typically support atomic manipulation of a single aggregate at a time

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Four Basic Types of NoSQL Databases

- Key-value stores
- Document databases
- Column-family stores
- Graph databases

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Key-value Stores: Representatives

LevelDB

Project Voldemort

Zdenka Prokopova
TBU in Zlín

Ranked list: http://db-engines.com/en/ranking/key-value+store

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Document Databases: Basics

- Basic concept of data: *Document*
- Documents are self-describing pieces of data
  - Hierarchical tree data structures
  - Nested associative arrays (maps), collections, scalars
  - XML, JSON (JavaScript Object Notation), BSON, …
- Documents in a collection should be “similar”
  - Their schema can differ
- Documents stored in the value part of key-value
  - Key-value stores where the values are examinable
  - Building search indexes on various keys/fields

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Document Databases: Representatives

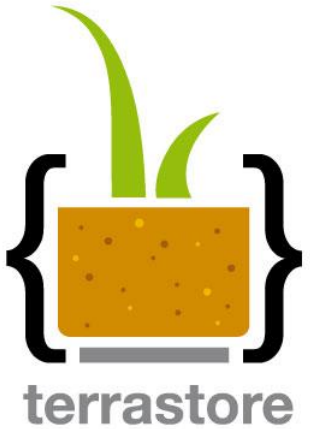

Zdenka Prokopova
TBU in Zlín

Ranked list: http://db-engines.com/en/ranking/document+store

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Column-family Stores: Basics

- wide-column, columnar
- Data model: rows that have many columns associated with a row key
- Column families are groups of related data (columns) that are often accessed together
  - e.g., for a customer we typically access all profile information at the same time, but not customer's orders

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Column-family Stores: Representatives

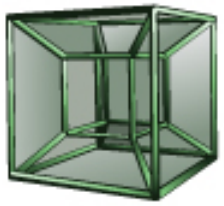

HYPERTABLE

accumulo™

Zdenka Prokopova
TBU in Zlín

Ranked list: http://db-engines.com/en/ranking/wide+column+store

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Graph Databases

- To store entities and relationships between them
  - Nodes are instances of objects
  - Nodes have properties, e.g., name
  - Edges have directional significance
  - Edges have types e.g., likes, friend, …

- Nodes are organized by relationships
  - Allow to find interesting patterns
  - example: Get all nodes that are “employee” of “Big Company” and that “likes” “NoSQL Distilled”

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Graph Databases: Example

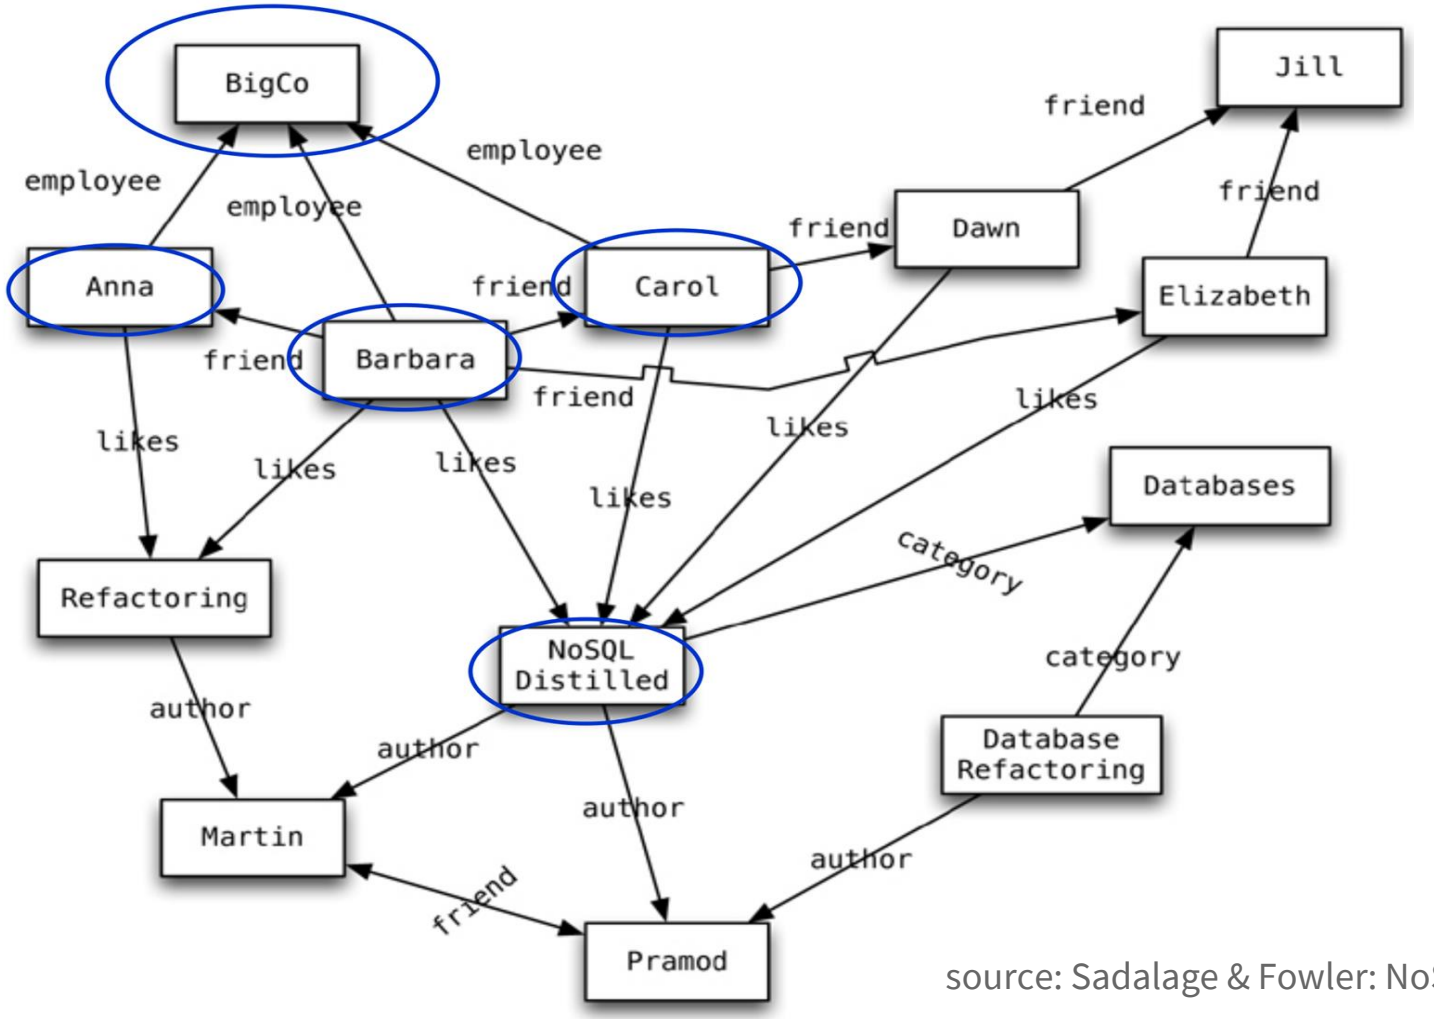

source: Sadalage & Fowler: NoSQL Distilled, 2012

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Graph Databases: Representatives

Zdenka Prokopova
TBU in Zlín

Ranked list: http://db-engines.com/en/ranking/graph+dbms

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION
YOUTH AND SPORTS

One Example of NoSQL Usage

facebook

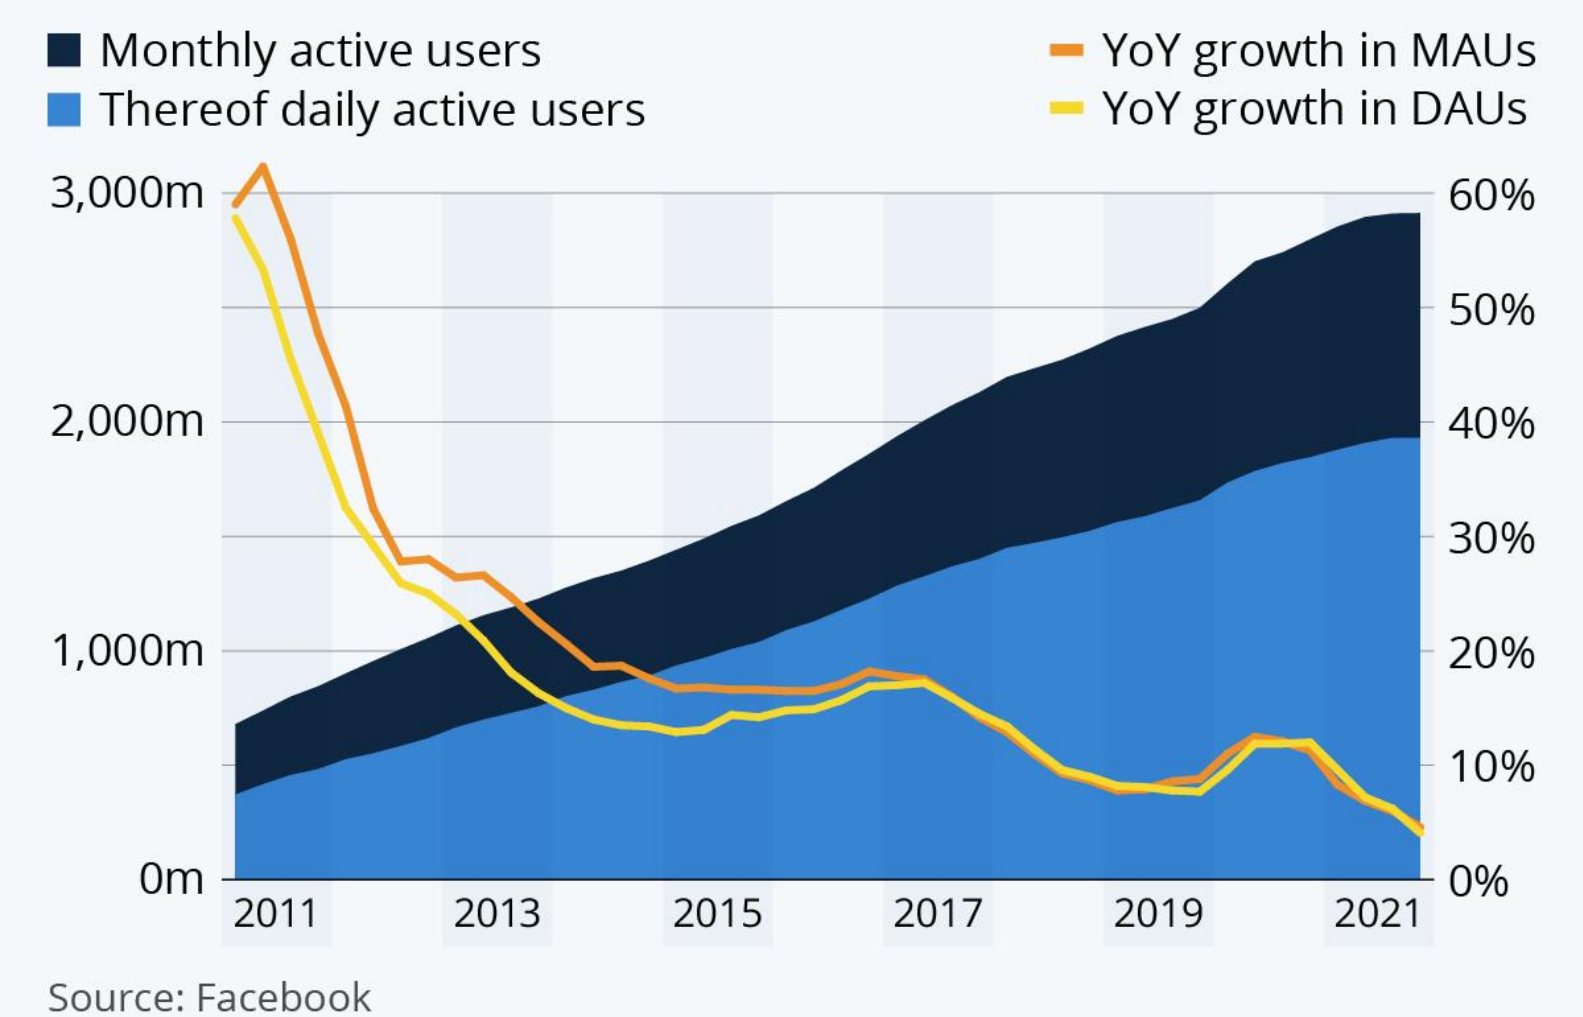

Zdenka Prokopova
TBU in Zlín

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Facebook: Database Technology Behind

Apache Hadoop http://hadoop.apache.org/

- Hadoop File System (HDFS)
  - over 100 PB in a single HDFS cluster
- an open source implementation of MapReduce:
  - Enables efficient calculations on massive amounts of data

Apache Hive http://hive.apache.org/

- SQL-like access to Hadoop-stored data
- integration of MapReduce query evaluation

Zdenka Prokopova
TBU in Zlín

sources: http://goo.gl/SZ6jia http://royal.pingdom.com/2010/06/18/the-software-behind-facebook/

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Facebook: Database Technology Behind

Apache HBase http://hbase.apache.org/

- a Hadoop column-family database
- used for e-mails, instant messaging and SMS
- replacement for MySQL and Cassandra

Memcached http://memcached.org/

- distributed key-value store
- used as a cache between web servers and MySQL servers since the beginning of FB

Zdenka Prokopova
TBU in Zlín

sources: http://goo.gl/SZ6jia http://royal.pingdom.com/2010/06/18/the-software-behind-facebook/

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Facebook: Database Technology Behind

Apache Giraph http://giraph.apache.org/

- graph database
- facebook users and connections is one very large graph
- used since 2013 for various analytic tasks

RocksDB http://rocksdb.org/

- high-performance key-value store
- developed internally in FB, now open-source

Zdenka Prokopova
TBU in Zlín

sources: https://code.facebook.com/posts/509727595776839/scaling-apache-giraph-to-a-trillion-edges/ http://goo.gl/XNtG6p

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

DB-Engines Ranking

Rank			DBMS	Database Model
Feb 2022	Jan 2022	Feb 2021		
1.	1.	1.	Oracle	Relational, Multi-model
2.	2.	2.	MySQL	Relational, Multi-model
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model
4.	4.	4.	PostgreSQL	Relational, Multi-model
5.	5.	5.	MongoDB	Document, Multi-model
6.	6.	↑ 7.	Redis	Key-value, Multi-model
7.	7.	↓ 6.	IBM Db2	Relational, Multi-model
8.	8.	8.	Elasticsearch	Search engine, Multi-model
9.	9.	↑ 11.	Microsoft Access	Relational
10.	10.	↓ 9.	SQLite	Relational
11.	11.	↓ 10.	Cassandra	Wide column
12.	12.	12.	MariaDB	Relational, Multi-model

Zdenka Prokopova
TBU in Zlín

DB-Engines Ranking. (2022). Retrieved from https://db-engines.com/en/ranking

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

DBMS popularity per category

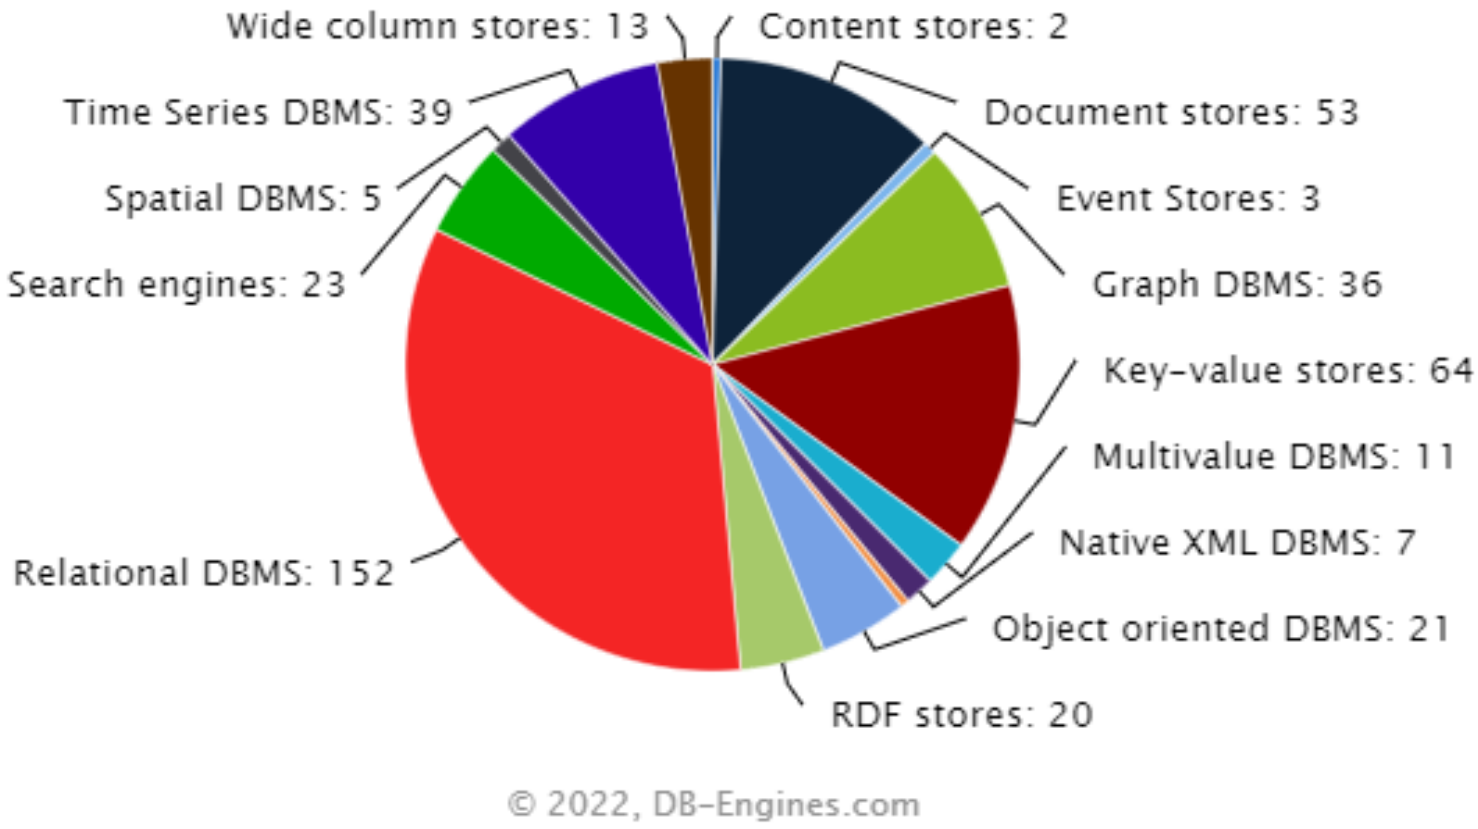

Zdenka Prokopova
TBU in Zlín

DB-Engines Ranking. (2022). Retrieved from https://db-engines.com/en/ranking

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

References

- Erl, T., Khattak, W., & Buhler, P. (2016). *Big Data Fundamentals: Concepts, Drivers \& Techniques*: Prentice Hall Press.
- Sadalage, P. J., & Fowler, M. (2013). *NoSQL distilled: a brief guide to the emerging world of polyglot persistence*: Pearson Education.
- Strauch, C., Sites, U.-L. S., & Kriha, W. (2011). NoSQL databases. *Lecture Notes, Stuttgart Media University, 20*.
- DB-Engines Ranking. (2019). Retrieved from https://db-engines.com/en/ranking
- Gain the insights, advice and tools (2019). Retrieved from https://www.gartner.com/en

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

References

- *RNDr. Irena Holubova, Ph.D. MMF UK course PA195: NoSQL Databases*
- *Data science and big data analytics*. (2018). New York, NY: Springer Berlin Heidelberg.
- Deka, G. C. (2017). *NoSQL : database for storage and retrieval of data in cloud*. Boca Raton, FL: CRC Press, Taylor & Francis Group.
- Perkins, L., Redmond, E., & Wilson, J. R. (2018). *Seven databases in seven weeks : a guide to modern databases and the NoSQL movement* (Second edition. ed.). Raleigh, North Carolina: The Pragmatic Bookshelf.
- Wiese, L. (2015). *Advanced data management : for SQL, NoSQL, cloud and distributed databases*. Berlin ; Boston: De Gruyter, Oldenbourg.

EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education

MINISTRY OF EDUCATION,
YOUTH AND SPORTS

Questions?

Strategic project of TBU in Zlín, reg. no. CZ.02.2.69/0.0/0.0/16_015/0002204

Zdenka Prokopova
TBU in Zlín